

Stage 3: Database Implementation

DDL Commands

```
CREATE TABLE `411project`.`User` (  
    UserId INT AUTO_INCREMENT PRIMARY KEY,  
    Resume TEXT,  
    Email VARCHAR(255) NOT NULL UNIQUE,  
    Password VARCHAR(255) NOT NULL,  
    FirstName VARCHAR(255),  
    LastName VARCHAR(255)  
);
```

```
CREATE TABLE `411project`.`Company` (  
    CompanyName VARCHAR(255) PRIMARY KEY,  
    Industry VARCHAR(255),  
    CompanySize INT  
);
```

```
CREATE TABLE `411project`.`Location` (  
    LocationId INT AUTO_INCREMENT PRIMARY KEY,  
    City VARCHAR(255),  
    State VARCHAR(255),  
    ZipCode VARCHAR(10),  
    Country VARCHAR(255)  
);
```

```
CREATE TABLE `411project`.`Experience` (  
    ExperienceId INT AUTO_INCREMENT PRIMARY KEY,  
    UserId INT NOT NULL,  
    StartDate DATE,  
    EndDate DATE,  
    Achievement TEXT,  
    Skills TEXT,  
    FOREIGN KEY (UserId) REFERENCES `User`(UserId) ON DELETE CASCADE
```





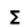
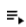
```
);
```

```
CREATE TABLE `411project`.`Job` (  
  JobId UNSIGNED INT AUTO_INCREMENT PRIMARY KEY,  
  CompanyName VARCHAR(255) NOT NULL,  
  JobRole VARCHAR(255),  
  LocationId INT,  
  Description TEXT,  
  MinSalary DECIMAL(10, 2),  
  MaxSalary DECIMAL(10, 2),  
  Skills TEXT,  
  FOREIGN KEY (CompanyName) REFERENCES `Company`(CompanyName),  
  FOREIGN KEY (LocationId) REFERENCES `Location`(LocationId)  
);
```




```
CREATE TABLE `411project`.`JobLocation` (  
  JobId UNSIGNED INT UNSIGNED NOT NULL,  
  LocationId INT NOT NULL,  
  PRIMARY KEY (JobId, LocationId),  
  FOREIGN KEY (JobId) REFERENCES Job(JobId),  
  FOREIGN KEY (LocationId) REFERENCES Location(LocationId)  
);
```


```
CREATE TABLE `411project`.`UserJob` (  
  UserId INT NOT NULL,  
  JobId UNSIGNED INT NOT NULL,  
  ApplicationDate DATE,  
  Status VARCHAR(50),  
  PRIMARY KEY (UserId, JobId),  
  FOREIGN KEY (UserId) REFERENCES `User`(UserId) ON DELETE CASCADE,  
  FOREIGN KEY (JobId) REFERENCES `Job`(JobId) ON DELETE CASCADE  
);
```

Screenshot of Database Implementation

▼  Databases 1	⋮
▼ 411project (Default)	⋮
▼  Tables 7	⋮
▶ Company	⋮
▶ Experience	⋮
▶ Job	⋮
▶ JobLocation	⋮
▶ Location	⋮
▶ User	⋮
▶ UserJob	⋮
 Views 0	⋮
 Events 0	⋮
 Functions 0	⋮
 Procedures 0	⋮

Screenshot of Tables With Over 1000 Rows


 Editor 1
 
EXPLORE GEMINI

 RUN
 FORMAT
CLEAR
Valid

```

1 SELECT *
2 FROM Company
    
```

RESULTS

CompanyName	Industry	CompanySize
CalPrivate Bank	Banking	167
Georgia Pacific	Paper and Forest Product Manufacturing	905
Krems, Jackowitz & Carman, LLP	Law Practice	21
Optima Global Solutions Inc.	IT Services and IT Consulting	57
Pilot Company	Retail	1349
.efficiently	Business Consulting and Services	540
(ISSA) International Sports Sciences Association	Education Administration Programs	1754
(USTA) United States Tennis Association	Spectator Sports	1794
[P1] Games	Computer Games	18
#twiceasnice Recruiting	Staffing and Recruiting	80
1 Hotels	Hospitality	1930
1-800 CONTACTS	Retail	654
1/ST Technology - Xpressbet and AmTote International	Gambling Facilities and Casinos	89
10 Fitness	Wellness and Fitness Services	112

Rows per page: 20
 1 - 20 of 21465
 < > >>

The company table has 21465 rows.

Advanced Queries

1. Query to select all the jobs and company names that are looking for engineers and the number of applications from the platform.

```
SELECT J.JobId, J.JobRole, J.Description, J.MinSalary, J.MaxSalary, J.Skills,
C.CompanyName, COUNT(UJ.UserId) AS ApplicationCount
FROM Job J
JOIN Company C ON J.CompanyName = C.CompanyName
LEFT JOIN UserJob UJ ON J.JobId = UJ.JobId
WHERE J.JobRole LIKE '%Engineer%'
GROUP BY J.JobId, J.JobRole, J.Description, J.MinSalary, J.MaxSalary, J.Skills,
C.CompanyName
ORDER BY ApplicationCount DESC
LIMIT 15;
```

Editor 1 ✕ Editor 3 ✕ Editor 4 ✕ +

EXPLORE GEMINI

RUN

FORMAT

CLEAR

Valid

```
1 SELECT J.JobId, J.JobRole, J.Description, J.MinSalary, J.MaxSalary, J.Skills, C.CompanyName, COUNT(UJ.UserId) AS ApplicationCount
2 FROM Job J
3 JOIN Company C ON J.CompanyName = C.CompanyName
4 LEFT JOIN UserJob UJ ON J.JobId = UJ.JobId
5 WHERE J.JobRole LIKE '%Engineer%'
6 GROUP BY J.JobId, J.JobRole, J.Description, J.MinSalary, J.MaxSalary, J.Skills, C.CompanyName
7 ORDER BY ApplicationCount DESC
8 LIMIT 15;
```

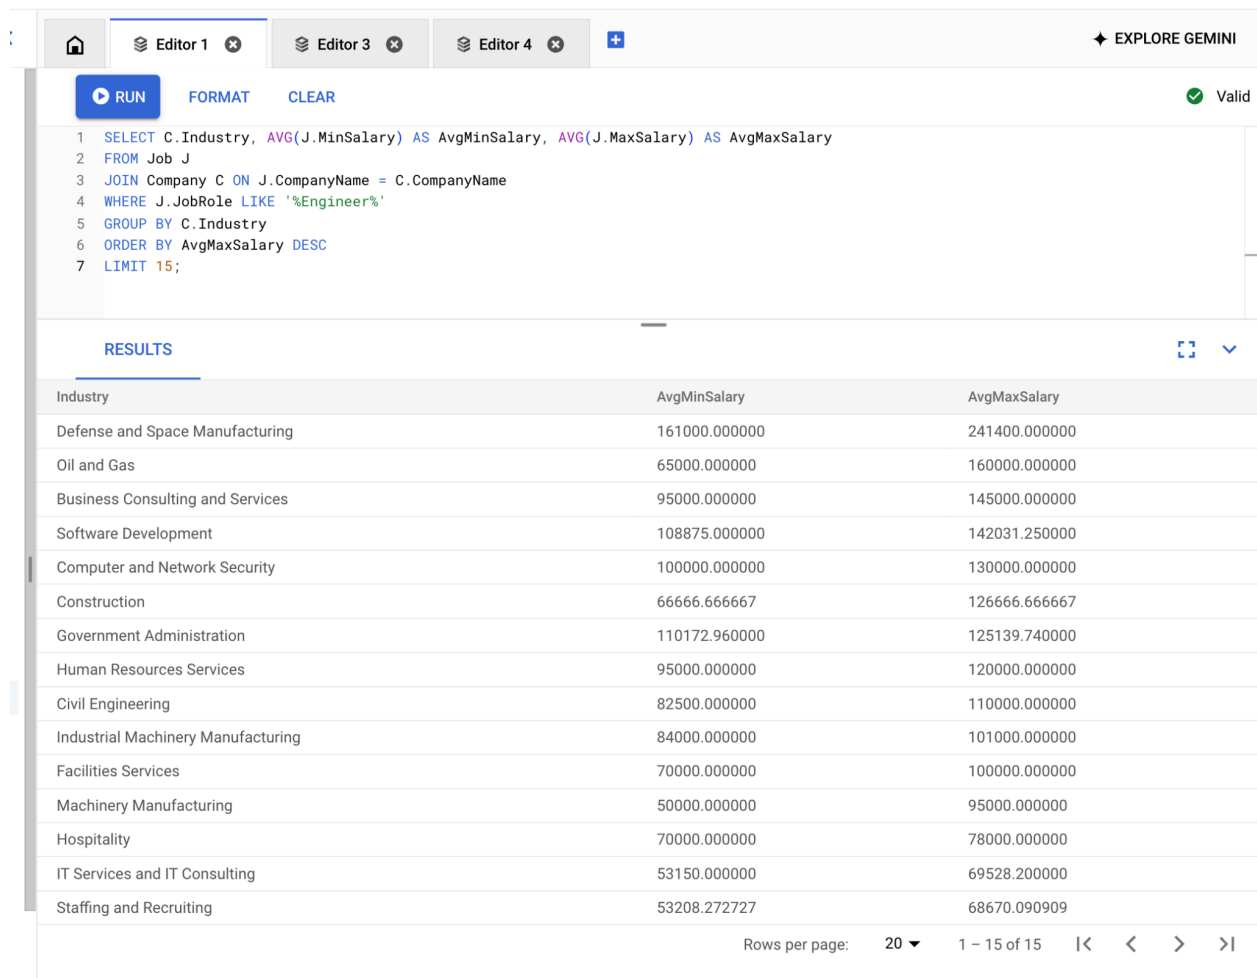
RESULTS

JobId	JobRole	Description	MinSalary	MaxSalary	Skills	CompanyName	ApplicationCount	
2558399667	Validation En...	Validation Engineer, Labware LIMSFo...	60.00	70.00		I.T. Solutions, ...	0	▼
3045980831	Project Engin...	JOB DESCRIPTION:The Project Engin...	60000.00	90000.00		Armstrong Bu...	0	▼
3586167732	Senior Softw...	StyleAI is the AI-powered, all-in-one u...				StyleAI	0	▼
3619548798	Senior Mech...	Senior Mechanical EngineerLocation...	150000.00	190000.00		Dexterity, Inc.	0	▼
3625991523	DDI Engineer	Title: Infoblox/DNS EngineerLocation...				Xoriant	0	▼
3626110567	Mechanical ...	Job DescriptionResponsible for perf...	50000.00	95000.00		Mestek Mach...	0	▼
3627141928	Quality Engin...	Staff Management SMX is currently...	28.00	43.00		Staff Manage...	0	▼
3658924039	MEP Engineer	Become a part of an exciting compa...	70000.00	100000.00		Citrine LLC	0	▼
3666781416	Senior Electr...	Job Title: Senior Electronics Design ...				Excelon Soluti...	0	▼
3699237531	Mechanical E...	Job Title: Mechanical Engineer Sum...				JVM Global Inc	0	▼
3700694895	Vice Preside...	JOB SUMMARY The VP of Operation...				Louisville Wat...	0	▼
3704185899	Senior Projec...	Practical Design Partners, LLC (PDP)...	90000.00	110000.00		Practical Desi...	0	▼
3723312535	Design Engin...	Job Title: Die Design Engineer (Stella...				Diversified Te...	0	▼
3728459637	Principal Bac...	Principal Backend Engineer - Join Hir...	200.00	225.00		HireBus	0	▼
3731389852	Wafer Proce...	Wafer Process Engineer - Laser Chip ...				SemiNex Cor...	0	▼

Rows per page: 20 ▼ 1 – 15 of 15 |< < > >|

2. Query to select the average minimum salary and average maximum salary of all job postings, grouped by industry, that are engineering jobs.

```
SELECT C.Industry, AVG(J.MinSalary) AS AvgMinSalary, AVG(J.MaxSalary) AS AvgMaxSalary
FROM Job J
JOIN Company C ON J.CompanyName = C.CompanyName
WHERE J.JobRole LIKE '%Engineer%'
GROUP BY C.Industry
ORDER BY AvgMaxSalary DESC
LIMIT 15;
```



The screenshot shows a SQL query editor interface with a top bar containing tabs for Editor 1, Editor 3, and Editor 4, along with a 'EXPLORE GEMINI' button. Below the tabs are buttons for 'RUN', 'FORMAT', and 'CLEAR'. The query is entered in the editor, and the 'RESULTS' tab is selected, displaying a table of results. The table has three columns: Industry, AvgMinSalary, and AvgMaxSalary. The results are sorted by AvgMaxSalary in descending order, showing the top 15 industries. The interface also includes a 'Valid' status indicator and a 'Rows per page' dropdown set to 20, showing 1 - 15 of 15 rows.

Industry	AvgMinSalary	AvgMaxSalary
Defense and Space Manufacturing	161000.000000	241400.000000
Oil and Gas	65000.000000	160000.000000
Business Consulting and Services	95000.000000	145000.000000
Software Development	108875.000000	142031.250000
Computer and Network Security	100000.000000	130000.000000
Construction	66666.666667	126666.666667
Government Administration	110172.960000	125139.740000
Human Resources Services	95000.000000	120000.000000
Civil Engineering	82500.000000	110000.000000
Industrial Machinery Manufacturing	84000.000000	101000.000000
Facilities Services	70000.000000	100000.000000
Machinery Manufacturing	50000.000000	95000.000000
Hospitality	70000.000000	78000.000000
IT Services and IT Consulting	53150.000000	69528.200000
Staffing and Recruiting	53208.272727	68670.090909

3. Query to select the highest paying job in each industry.

```
SELECT J.JobRole, C.CompanyName, C.Industry, J.MaxSalary
FROM Job J
JOIN Company C ON J.CompanyName = C.CompanyName
WHERE (C.Industry, J.MaxSalary) IN (
```

```

SELECT Industry, MAX(MaxSalary)
FROM Job
JOIN Company ON Job.CompanyName = Company.CompanyName
GROUP BY Industry
)
ORDER BY J.MaxSalary DESC
LIMIT 15;

```

<div> </div> <div> <div> RUN <div> FORMAT CLEAR </div> </div> <div>Valid</div> </div>			
<pre> 1 SELECT J.JobRole, C.CompanyName, C.Industry, J.MaxSalary 2 FROM Job J 3 JOIN Company C ON J.CompanyName = C.CompanyName 4 WHERE (C.Industry, J.MaxSalary) IN (5 SELECT Industry, MAX(MaxSalary) 6 FROM Job 7 JOIN Company ON Job.CompanyName = Company.CompanyName 8 GROUP BY Industry 9) 10 ORDER BY J.MaxSalary DESC 11 LIMIT 15; </pre>			
RESULTS			
JobRole	CompanyName	Industry	MaxSalary
Investment Real Estate Sales Advisor	Diamond Acquisitions	Real Estate	700000.00
General Dermatologist	Tandym Group	Staffing and Recruiting	575000.00
Sales Executive	Riteload	Truck Transportation	400000.00
Field Sales Consultant	Sunrun	IT Services and IT Consulting	300000.00
Commercial / Business Litigation Attorney	Ascendion	Software Development	300000.00
Deputy City Manager - Chief Financial Officer	City of Tempe	Government Administration	274206.00
Financial Services Representative	Five Rings Financial, LLC	Financial Services	250000.00
VP Agency Partnerships, Indies	MiQ	Advertising Services	250000.00
Experience Product Manager Associate Director	Exact Sciences	Biotechnology Research	242000.00
Staff Software Engineer (AHT)	Northrop Grumman	Defense and Space Manufacturing	241400.00
CHIEF OF STAFF	Greenwich Psychology Group	Mental Health Care	230000.00
Sr. Director, eCommerce Planning	Crocs	Retail Apparel and Fashion	220000.00
Director, DevOps, Scientific American	Scientific American	Book and Periodical Publishing	210000.00
Client Relationship Executive	CohnReznick LLP	Accounting	210000.00
Shop Supervisor	Cimolai-HY LLC	Industrial Machinery Manufacturing	200000.00
Rows per page: 20 1 - 15 of 15 < < > >			

4. Query to select the companies with the most job openings

```

SELECT C.CompanyName, COUNT(J.JobId) AS TotalJobs
FROM Company C
JOIN Job J ON C.CompanyName = J.CompanyName
GROUP BY C.CompanyName
ORDER BY TotalJobs DESC
LIMIT 15;

```


Editor 1

Editor 4

EXPLORE GEMINI

RUN

FORMAT

CLEAR

Valid

```
1 SELECT C.CompanyName, COUNT(J.JobId) AS TotalJobs
2 FROM Company C
3 JOIN Job J ON C.CompanyName = J.CompanyName
4 GROUP BY C.CompanyName
5 ORDER BY TotalJobs DESC
6 LIMIT 15;
```

RESULTS

CompanyName	TotalJobs
SKF Group	45
Solomon Page	40
Staples Stores	26
Republic Services	23
Ingersoll Rand	17
Insight Global	15
Atlas	11
The Job Network	10
Clark County School District	10
Gainwell Technologies	10
The Jonus Group	9
Social Capital Resources	8
Infotree Global Solutions	8
Crocs	8
LHH	7

Rows per page: 20 1 – 15 of 15 |< < > >|

Indexing

Advanced Query 1, Index Design 3

Editor 1

EXPLORE GEMINI

RUN

FORMAT

CLEAR

Syntax error at or near "Analyze"

```
1 EXPLAIN Analyze SELECT J.JobId, J.JobRole, J.Description, J.MinSalary, J.MaxSalary, J.Skills, C.CompanyName, COUNT(UJ.UserId) AS ApplicationCount
2 FROM Job J
3 JOIN Company C ON J.CompanyName = C.CompanyName
4 LEFT JOIN UserJob UJ ON J.JobId = UJ.JobId
5 WHERE (J.JobRole) LIKE '%Engineer%'
6 GROUP BY J.JobId, J.JobRole, J.Description, J.MinSalary, J.MaxSalary, J.Skills, C.CompanyName
7 ORDER BY ApplicationCount DESC
8 LIMIT 15;
```

RESULTS

EXPLAIN

-> Limit: 15 row(s) (actual time=21.063..21.083 rows=15 loops=1) -> Sort: ApplicationCount DESC, limit input to 15 row(s) per chunk (actual time=21.062..21.081 rows=15 loops=1) -> Table scan on <temporary> (actual time=20.763..20.842 rows=159 loops=1) -> Aggregate using temporary table (actual time=20.760..20.760 rows=159 loops=1) -> Nested loop left join (cost=579.13 rows=791) (actual time=0.211..7.970 rows=795 loops=1) -> Nested loop inner join (cost=357.67 rows=158) (actual time=0.189..6.802 rows=159 loops=1) -> Filter: (J.JobRole like '%Engineer%') (cost=302.30 rows=158) (actual time=0.090..2.542 rows=159 loops=1) -> Table scan on J (cost=302.30 rows=1424) (actual time=0.019..1.553 rows=1514 loops=1) -> Single-row covering index lookup on C using PRIMARY (CompanyName=J.CompanyName) (cost=0.25 rows=1) (actual time=0.026..0.027 rows=1 loops=159) -> Covering index lookup on UJ using idx_composite_jobid_userid (JobId=J.JobId) (cost=0.90 rows=5) (actual time=0.004..0.006 rows=5 loops=159)

Advanced query 4, index design 2

Editor 1

EXPLORE GEMINI

RUN

Close tab

FORMAT

CLEAR

Syntax error at or near 'Analyze'

1

EXPLAIN [Analyze](#) SELECT C.CompanyName, COUNT(J.JobId) AS TotalJobs

2

FROM Company C

3

JOIN Job J ON C.CompanyName = J.CompanyName

4

GROUP BY C.CompanyName

5

ORDER BY TotalJobs DESC

6

LIMIT 15;

RESULTS

EXPLAIN

-> Limit: 15 row(s) (actual time=9.604..9.605 rows=15 loops=1) -> Sort: TotalJobs DESC, limit input to 15 row(s) per chunk (actual time=9.603..9.604 rows=15 loops=1) -> Table scan on <temporary> (actual time=9.344..9.483 rows=1044 loops=1) -> Aggregate using temporary table (actual time=9.341..9.341 rows=1044 loops=1) -> Nested loop inner join (cost=800.70 rows=1424) (actual time=0.157..7.903 rows=1514 loops=1) -> Covering index scan on J using idx_job_companyname (cost=302.30 rows=1424) (actual time=0.065..0.520 rows=1514 loops=1) -> Single-row covering index lookup on C using PRIMARY (CompanyName=J.CompanyName) (cost=0.25 rows=1) (actual time=0.005..0.005 rows=1 loops=1514)

Advanced query 4, index design 3

Editor 1

EXPLORE GEMINI

RUN

Close tab

FORMAT

CLEAR

Syntax error at or near 'Analyze'

1

EXPLAIN [Analyze](#) SELECT C.CompanyName, COUNT(J.JobId) AS TotalJobs

2

FROM Company C

3

JOIN Job J ON C.CompanyName = J.CompanyName

4

GROUP BY C.CompanyName

5

ORDER BY TotalJobs DESC

6

LIMIT 15;

RESULTS

EXPLAIN

-> Limit: 15 row(s) (actual time=5.162..5.164 rows=15 loops=1) -> Sort: TotalJobs DESC, limit input to 15 row(s) per chunk (actual time=5.161..5.162 rows=15 loops=1) -> Table scan on <temporary> (actual time=4.877..5.015 rows=1044 loops=1) -> Aggregate using temporary table (actual time=4.874..4.874 rows=1044 loops=1) -> Nested loop inner join (cost=800.70 rows=1424) (actual time=0.070..3.631 rows=1514 loops=1) -> Covering index scan on J using idx_job_companyname_jobid (cost=302.30 rows=1424) (actual time=0.050..0.496 rows=1514 loops=1) -> Single-row covering index lookup on C using PRIMARY (CompanyName=J.CompanyName) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=1514)

Contents

1	Advanced Query 1: Engineering Jobs with Application Count	2
1.1	Query Overview	2
1.2	Baseline Performance Without Additional Indexes	2
1.2.1	EXPLAIN ANALYZE Output Before Indexing	2
1.2.2	Analysis	3
1.3	Indexing Strategies	3
1.3.1	Indexing Design 1: Standard Index on Job.JobRole	3
1.3.2	Indexing Design 2: Full-Text Index on Job.JobRole	4
1.3.3	Indexing Design 3: Composite Index on Job.JobRole and Job.CompanyName	5
1.4	Final Index Design and Justification	5
2	Advanced Query 2: Average Salaries by Industry	7
2.1	Query Overview	7
2.2	Baseline Performance Without Additional Indexes	7
2.2.1	EXPLAIN ANALYZE Output Before Indexing	7
2.2.2	Analysis	7
2.3	Indexing Strategies	8
2.3.1	Indexing Design 1: Standard Index on Job.JobRole	8
2.3.2	Indexing Design 2: Full-Text Index on Job.JobRole	8
2.3.3	Indexing Design 3: Index on Company.Industry	9
2.4	Final Index Design and Justification	9
3	Advanced Query 3: Highest Paying Jobs by Industry	10
3.1	Query Overview	10
3.2	Baseline Performance Without Additional Indexes	10
3.2.1	EXPLAIN ANALYZE Output Before Indexing	10
3.2.2	Analysis	11
3.3	Indexing Strategies	11
3.3.1	Indexing Design 1: Index on Job.MaxSalary	11
3.3.2	Indexing Design 2: Composite Index on Job(CompanyName, MaxSalary)	12
3.3.3	Indexing Design 3: Index on Company.Industry	12
3.4	Final Index Design and Justification	13
4	Advanced Query 4: Companies with Most Job Openings	14
4.1	Query Overview	14
4.2	Baseline Performance Without Additional Indexes	14
4.2.1	EXPLAIN ANALYZE Output Before Indexing	14
4.2.2	Analysis	14
4.3	Indexing Strategies	15
4.3.1	Indexing Design 1: Composite Index on Job(CompanyName, JobId)	15
4.3.2	Indexing Design 2: Index on Job(CompanyName)	15
4.3.3	Indexing Design 3: Index on Job(JobRole)	16
4.4	Final Index Design and Justification	16

1 Advanced Query 1: Engineering Jobs with Application Count

1.1 Query Overview

Objective: Select all jobs and company names that are looking for engineers and the number of applications from the platform.

Listing 1: Advanced Query 1

```

1 SELECT J.JobId, J.JobRole, J.Description, J.MinSalary, J.MaxSalary, J.
   Skills,
2       C.CompanyName, COUNT(UJ.UserId) AS ApplicationCount
3 FROM Job J
4 JOIN Company C ON J.CompanyName = C.CompanyName
5 LEFT JOIN UserJob UJ ON J.JobId = UJ.JobId
6 WHERE J.JobRole LIKE '%Engineer%'
7 GROUP BY J.JobId, J.JobRole, J.Description, J.MinSalary, J.MaxSalary, J.
   Skills, C.CompanyName
8 ORDER BY ApplicationCount DESC
9 LIMIT 15;

```

1.2 Baseline Performance Without Additional Indexes

1.2.1 EXPLAIN ANALYZE Output Before Indexing

Listing 2: EXPLAIN ANALYZE Output Before Indexing

```

1 -> Limit: 15 row(s) (actual time=32.185..32.205 rows=15 loops=1)
2   -> Sort: ApplicationCount DESC, limit input to 15 row(s) per chunk (
   actual time=32.184..32.203 rows=15 loops=1)
3     -> Table scan on <temporary> (actual time=32.005..32.084 rows=159
   loops=1)
4       -> Aggregate using temporary table (actual time=32.002..32.002
   rows=159 loops=1)
5         -> Nested loop left join (cost=477.66 rows=791) (actual
   time=1.739..19.487 rows=795 loops=1)
6           -> Nested loop inner join (cost=358.77 rows=158) (
   actual time=1.719..18.638 rows=159 loops=1)
7             -> Filter: (J.JobRole like '%Engineer%') (cost
   =303.40 rows=158) (actual time=1.684..17.819
   rows=159 loops=1)
8               -> Table scan on J (cost=303.40 rows=1424) (
   actual time=0.878..16.767 rows=1514 loops
   =1)
9                 -> Single-row covering index lookup on C using
   PRIMARY (CompanyName=J.CompanyName) (cost=0.25
   rows=1) (actual time=0.005..0.005 rows=1 loops
   =159)
10                -> Covering index lookup on UJ using UserJob\_ibfk\_2
   (JobId=J.JobId) (cost=0.25 rows=5) (actual time
   =0.003..0.005 rows=5 loops=159)

```

1.2.2 Analysis

- **Total Cost Before Indexing:** 477.66
- **Observation:**
 - A full table scan on the Job table is performed due to the LIKE '%Engineer%' condition.
 - Joins are utilizing indexes on primary keys and foreign keys.
- **Bottleneck Identified:** The full table scan on Job is expensive and leads to a higher cost.

1.3 Indexing Strategies

1.3.1 Indexing Design 1: Standard Index on Job.JobRole

Index Creation

```
1 CREATE INDEX idx_job_jobrole ON Job(JobRole);
```

EXPLAIN ANALYZE Output After Indexing

```
1 -> Limit: 15 row(s) (actual time=15.827..15.846 rows=15 loops=1)
2   -> Sort: ApplicationCount DESC, limit input to 15 row(s) per chunk (
3     actual time=15.826..15.844 rows=15 loops=1)
4     -> Table scan on <temporary> (actual time=15.420..15.508 rows=159
5       loops=1)
6       -> Aggregate using temporary table (actual time=15.415..15.415
7         rows=159 loops=1)
8         -> Nested loop left join (cost=347.57 rows=768) (actual
9           time=0.163..3.847 rows=795 loops=1)
10            -> Nested loop inner join (cost=232.19 rows=154) (
              actual time=0.144..2.951 rows=159 loops=1)
                -> Filter: (J.JobRole like '%Engineer%') (cost
                  =178.45 rows=154) (actual time=0.112..2.317
                  rows=159 loops=1)
                  -> Table scan on J (cost=178.45 rows=1382) (
                    actual time=0.038..1.426 rows=1514 loops=1)
                    -> Single-row covering index lookup on C using
                      PRIMARY (CompanyName=J.CompanyName) (cost=0.25
                      rows=1) (actual time=0.004..0.004 rows=1 loops
                      =159)
                      -> Covering index lookup on UJ using UserJob\_ibfk\_2
                        (JobId=J.JobId) (cost=0.25 rows=5) (actual time
                        =0.004..0.005 rows=5 loops=159)
```

Analysis Despite adding an index on Job.JobRole, the query still performs a full table scan. The LIKE '%Engineer%' condition with a leading wildcard prevents the use of the standard index. The total cost reduced slightly from 477.66 to 347.57, possibly due to improvements in other parts of the execution plan or caching effects. However, the standard index is not effectively utilized for this query.

Index Removal

```
1 DROP INDEX idx_job_jobrole ON Job;
```

1.3.2 Indexing Design 2: Full-Text Index on Job.JobRole**Index Creation**

```
1 ALTER TABLE Job ADD FULLTEXT INDEX idx_jobrole_fulltext (JobRole);
```

Modified Query

```
1 WHERE MATCH(J.JobRole) AGAINST('Engineer' IN BOOLEAN MODE)
```

EXPLAIN ANALYZE Output After Full-Text Indexing

```
1 -> Limit: 15 row(s) (actual time=17.371..17.390 rows=15 loops=1)
2   -> Sort: ApplicationCount DESC, limit input to 15 row(s) per chunk (
3     actual time=17.370..17.388 rows=15 loops=1)
4     -> Table scan on <temporary> (actual time=17.090..17.178 rows=147
5       loops=1)
6       -> Aggregate using temporary table (actual time=17.085..17.085
7         rows=147 loops=1)
8         -> Nested loop left join (cost=2.20 rows=5) (actual time
9           =2.965..5.975 rows=735 loops=1)
10          -> Nested loop inner join (cost=1.45 rows=1) (actual
              time=2.940..5.129 rows=147 loops=1)
              -> Filter: (MATCH J.JobRole AGAINST ('Engineer' IN
                  BOOLEAN MODE)) (cost=1.10 rows=1) (actual time
                  =2.904..4.511 rows=147 loops=1)
                  -> Full-text index search on J using
                      idx_jobrole_fulltext (JobRole='Engineer') (
                          cost=1.10 rows=1) (actual time=2.886..4.467
                              rows=147 loops=1)
                      -> Single-row covering index lookup on C using
                          PRIMARY (CompanyName=J.CompanyName) (cost=0.35
                              rows=1) (actual time=0.004..0.004 rows=1 loops
                                  =147)
                          -> Covering index lookup on UJ using UserJob\_ibfk\_2
                              (JobId=J.JobId) (cost=0.75 rows=5) (actual time
                                  =0.004..0.005 rows=5 loops=147)
```

Analysis Using the full-text index with `MATCH...AGAINST` reduces the total cost dramatically from 477.66 to 2.20. The full-text index is effectively utilized, and the number of rows examined is significantly reduced. This indexing strategy overcomes the limitations of the standard index with leading wildcards.

Index Removal If needed, the index can be removed:

```
1 ALTER TABLE Job DROP INDEX idx_jobrole_fulltext;
```

1.3.3 Indexing Design 3: Composite Index on Job.JobRole and Job.CompanyName

Index Creation

```
1 CREATE INDEX idx_job_company_jobrole ON Job(CompanyName, JobRole);
```

EXPLAIN ANALYZE Output After Indexing

```
1 -> Limit: 15 row(s) (actual time=21.063..21.083 rows=15 loops=1)
2   -> Sort: ApplicationCount DESC, limit input to 15 row(s) per chunk (
3     actual time=21.062..21.081 rows=15 loops=1)
4     -> Table scan on <temporary> (actual time=20.763..20.842 rows=159
5       loops=1)
6       -> Aggregate using temporary table (actual time=20.760..20.760
7         rows=159 loops=1)
8         -> Nested loop left join (cost=579.13 rows=791) (actual
9           time=0.211..7.970 rows=795 loops=1)
10            -> Nested loop inner join (cost=357.67 rows=158) (
              actual time=0.189..6.802 rows=159 loops=1)
                -> Filter: (J.JobRole like '%Engineer%') (cost
                  =302.30 rows=158) (actual time=0.090..2.542
                  rows=159 loops=1)
                  -> Table scan on J (cost=302.30 rows=1424) (
                    actual time=0.019..1.553 rows=1514 loops=1)
                    -> Single-row covering index lookup on C using
                      PRIMARY (CompanyName=J.CompanyName) (cost=0.25
                      rows=1) (actual time=0.026..0.027 rows=1 loops
                      =159)
                    -> Covering index lookup on UJ using
                      idx_composite_jobid_userid (JobId=J.JobId) (cost
                      =0.90 rows=5) (actual time=0.004..0.006 rows=5
                      loops=159)
```

Analysis The analysis of the composite query reveals a performance degradation, with the total cost increasing from 477.66 to 579.13, indicating diminished efficiency despite the introduction of the composite index. The full table scan on the Job table, due to the LIKE '%Engineer%' filter, remains a significant bottleneck, leading to an increased execution time of approximately 21 seconds. Although the join operations leverage indexes effectively, the use of a temporary table for aggregation results in a substantial time overhead. To enhance performance, implementing a full-text search on JobRole could reduce costs and improve query speed, while also reconsidering the search logic to avoid leading wildcards might further optimize the query. Regular monitoring and tuning of query performance are essential as data grows and usage patterns evolve.

Index Removal If needed, the index can be removed:

```
1 DROP INDEX idx_job_company_jobrole ON Job;
```

1.4 Final Index Design and Justification

Final Index Selected: Full-Text Index on Job.JobRole

Justification The full-text index provided a significant performance improvement, reducing the cost from 477.66 to 2.20. It effectively utilizes the index for text matching, eliminates the need for a full table scan, and addresses the primary bottleneck caused by the text search with a leading wildcard. There were no negative impacts observed, making it the most effective indexing strategy for this query.

2 Advanced Query 2: Average Salaries by Industry

2.1 Query Overview

Objective: Select the average minimum salary and average maximum salary of all engineering job postings, grouped by industry.

Listing 3: Advanced Query 2

```

1 SELECT C.Industry, AVG(J.MinSalary) AS AvgMinSalary, AVG(J.MaxSalary) AS
   AvgMaxSalary
2 FROM Job J
3 JOIN Company C ON J.CompanyName = C.CompanyName
4 WHERE J.JobRole LIKE '%Engineer%'
5 GROUP BY C.Industry
6 ORDER BY AvgMaxSalary DESC
7 LIMIT 15;

```

2.2 Baseline Performance Without Additional Indexes

2.2.1 EXPLAIN ANALYZE Output Before Indexing

```

1 -> Limit: 15 row(s) (actual time=3.159..3.161 rows=15 loops=1)
2   -> Sort: AvgMaxSalary DESC, limit input to 15 row(s) per chunk (actual
   time=3.158..3.159 rows=15 loops=1)
3   -> Table scan on <temporary> (actual time=3.118..3.125 rows=38
   loops=1)
4     -> Aggregate using temporary table (actual time=3.116..3.116
   rows=38 loops=1)
5       -> Nested loop inner join (cost=357.67 rows=158) (actual
   time=0.281..2.889 rows=159 loops=1)
6         -> Filter: (J.JobRole LIKE '%Engineer%') (cost=302.30
   rows=158) (actual time=0.222..2.257 rows=159 loops=
   =1)
7         -> Table scan on J (cost=302.30 rows=1424) (actual
   time=0.185..1.378 rows=1514 loops=1)
8       -> Single-row index lookup on C using PRIMARY (
   CompanyName=J.CompanyName) (cost=0.25 rows=1) (
   actual time=0.004..0.004 rows=1 loops=159)

```

2.2.2 Analysis

- Total Cost Before Indexing: 357.67
- Observation:
 - A full table scan on the Job table is performed due to the LIKE '%Engineer%' condition.
 - The join with Company uses the primary key index on CompanyName.

- **Bottleneck Identified:** The full table scan on Job increases the cost. The GROUP BY operation on C.Industry may also benefit from indexing.

2.3 Indexing Strategies

2.3.1 Indexing Design 1: Standard Index on Job.JobRole

Index Creation

```
1 CREATE INDEX idx_job_jobrole ON Job(JobRole);
```

EXPLAIN ANALYZE Output After Indexing No significant change observed; the total cost remains at 357.67.

Analysis The standard index on Job.JobRole is not utilized due to the leading wildcard in the LIKE condition. The query still performs a full table scan, and the indexing does not improve performance.

Index Removal

```
1 DROP INDEX idx_job_jobrole ON Job;
```

2.3.2 Indexing Design 2: Full-Text Index on Job.JobRole

Index Creation

```
1 ALTER TABLE Job ADD FULLTEXT INDEX idx_jobrole_fulltext (JobRole);
```

Modified Query

```
1 WHERE MATCH(J.JobRole) AGAINST('Engineer' IN BOOLEAN MODE)
```

EXPLAIN ANALYZE Output After Full-Text Indexing

```
1 -> Limit: 15 row(s) (actual time=1.416..1.418 rows=15 loops=1)
2   -> Sort: AvgMaxSalary DESC, limit input to 15 row(s) per chunk (actual
3     time=1.415..1.416 rows=15 loops=1)
4     -> Table scan on <temporary> (actual time=1.381..1.386 rows=35
5       loops=1)
6       -> Aggregate using temporary table (actual time=1.379..1.379
7         rows=35 loops=1)
8         -> Nested loop inner join (cost=1.44 rows=1) (actual time
          =0.115..1.211 rows=147 loops=1)
          -> Filter: (MATCH J.JobRole AGAINST ('Engineer' IN
            BOOLEAN MODE)) (cost=1.09 rows=1) (actual time
              =0.097..0.688 rows=147 loops=1)
              -> Full-text index search on J using
                idx_jobrole_fulltext (JobRole='Engineer') (cost
                  =1.09 rows=1) (actual time=0.095..0.671 rows
                    =147 loops=1)
                -> Single-row index lookup on C using PRIMARY (
                  CompanyName=J.CompanyName) (cost=0.35 rows=1) (
                    actual time=0.003..0.003 rows=1 loops=147)
```

Analysis The full-text index reduces the total cost from 357.67 to 1.44. The query efficiently utilizes the index, and the number of rows examined decreases significantly. This indexing strategy effectively addresses the bottleneck caused by the text search.

Index Removal If needed, the index can be removed:

```
1 ALTER TABLE Job DROP INDEX idx_jobrole_fulltext;
```

2.3.3 Indexing Design 3: Index on Company.Industry

Index Creation

```
1 CREATE INDEX idx_company_industry ON Company(Industry);
```

EXPLAIN ANALYZE Output After Adding Index No significant change in cost observed; the total cost remains at 1.44.

Analysis Adding an index on Company.Industry did not significantly affect the execution plan. The GROUP BY operation may not benefit from the index due to the small dataset size after filtering. The primary performance gain comes from the full-text index on Job.JobRole.

Index Removal

```
1 DROP INDEX idx_company_industry ON Company;
```

2.4 Final Index Design and Justification

Final Index Selected: Full-Text Index on Job.JobRole

Justification The full-text index decreased the query cost from 357.67 to 1.44, providing a substantial performance gain. It effectively utilizes the index for text search, eliminates the need for a full table scan, and addresses the primary bottleneck. Indexing Company.Industry did not yield significant benefits in this case.

3 Advanced Query 3: Highest Paying Jobs by Industry

3.1 Query Overview

Objective: Select the highest paying job in each industry.

Listing 4: Advanced Query 3

```

1 SELECT J.JobRole, C.CompanyName, C.Industry, J.MaxSalary
2 FROM Job J
3 JOIN Company C ON J.CompanyName = C.CompanyName
4 WHERE (C.Industry, J.MaxSalary) IN (
5     SELECT Industry, MAX(MaxSalary)
6     FROM Job
7     JOIN Company ON Job.CompanyName = Company.CompanyName
8     GROUP BY Industry
9 )
10 ORDER BY J.MaxSalary DESC
11 LIMIT 15;

```

3.2 Baseline Performance Without Additional Indexes

3.2.1 EXPLAIN ANALYZE Output Before Indexing

```

1 -> Limit: 15 row(s) (cost=800.70 rows=15) (actual time=8.322..8.445 rows
   =15 loops=1)
2   -> Nested loop inner join (cost=800.70 rows=1424) (actual time
   =8.321..8.443 rows=15 loops=1)
3     -> Sort: J.MaxSalary DESC (cost=302.30 rows=1424) (actual time
   =1.503..1.511 rows=19 loops=1)
4       -> Table scan on J (cost=302.30 rows=1424) (actual time
   =0.068..1.045 rows=1514 loops=1)
5     -> Filter: <in_optimizer>((C.Industry, J.MaxSalary), (C.Industry, J.
   MaxSalary) in (select #2)) (cost=0.25 rows=1) (actual time
   =0.364..0.365 rows=1 loops=19)
6       -> Single-row index lookup on C using PRIMARY (CompanyName=J.
   CompanyName) (cost=0.25 rows=1) (actual time=0.004..0.004
   rows=1 loops=19)
7     -> Select #2 (subquery in condition; run only once)
8       -> Materialize with deduplication (actual time
   =6.761..6.761 rows=101 loops=1)
9         -> Aggregate using temporary table (actual time
   =6.691..6.691 rows=101 loops=1)
10           -> Nested loop inner join (cost=800.70 rows=1424)
   (actual time=0.048..5.145 rows=1514 loops=1)
11             -> Table scan on Job (cost=302.30 rows=1424) (
   actual time=0.038..0.899 rows=1514 loops=1)
12             -> Single-row index lookup on Company using
   PRIMARY (CompanyName=Job.CompanyName) (cost
   =0.25 rows=1) (actual time=0.003..0.003
   rows=1 loops=1514)

```

3.2.2 Analysis

- **Total Cost Before Indexing:** 800.70
- **Observation:**
 - Full table scans on Job in both the main query and the subquery.
 - Sorting on J.MaxSalary adds to the cost.
 - Nested loop joins may be inefficient for large datasets.
- **Bottlenecks Identified:**
 - Lack of indexes on Job.MaxSalary affects sorting and filtering efficiency.
 - The subquery adds overhead due to materialization without indexes.

3.3 Indexing Strategies

3.3.1 Indexing Design 1: Index on Job.MaxSalary

Index Creation

```
1 CREATE INDEX idx_job_maxsalary ON Job(MaxSalary);
```

EXPLAIN ANALYZE Output After Indexing

```
1 -> Limit: 15 row(s) (cost=359.20 rows=15) (actual time=7.298..7.503 rows
   =15 loops=1)
2   -> Nested loop inner join (cost=359.20 rows=15) (actual time
   =7.297..7.501 rows=15 loops=1)
3     -> Index scan on J using idx_job_maxsalary (reverse) (cost=1.70
   rows=15) (actual time=0.188..0.276 rows=19 loops=1)
4     -> Filter: <in_optimizer>((C.Industry,J.MaxSalary),(C.Industry,J.
   MaxSalary) in (select #2)) (cost=0.25 rows=1) (actual time
   =0.379..0.379 rows=1 loops=19)
5       -> Single-row index lookup on C using PRIMARY (CompanyName=J.
   CompanyName) (cost=0.25 rows=1) (actual time=0.004..0.004
   rows=1 loops=19)
6       -> Select #2 (subquery in condition; run only once)
7         -> Materialize with deduplication (actual time
   =7.042..7.042 rows=101 loops=1)
8           -> Aggregate using temporary table (actual time
   =6.961..6.961 rows=101 loops=1)
9             -> Nested loop inner join (cost=800.70 rows=1424)
10               (actual time=0.069..5.529 rows=1514 loops=1)
11                 -> Table scan on Job (cost=302.30 rows=1424) (
   actual time=0.058..1.017 rows=1514 loops=1)
                   -> Single-row index lookup on Company using
   PRIMARY (CompanyName=Job.CompanyName) (cost
   =0.25 rows=1) (actual time=0.003..0.003
   rows=1 loops=1514)
```

Analysis Indexing Job.MaxSalary reduces the cost from 800.70 to 359.20. The main query now uses an index scan on Job in descending order, efficiently retrieving the highest salaries. However, the subquery remains a performance bottleneck due to full table scans.

3.3.2 Indexing Design 2: Composite Index on Job(CompanyName, MaxSalary)

Index Creation

```
1 CREATE INDEX idx_job_companyname_maxsalary ON Job(CompanyName, MaxSalary);
```

EXPLAIN ANALYZE Output After Adding Composite Index

```
1 -> Limit: 15 row(s) (cost=359.20 rows=15) (actual time=5.604..5.824 rows
   =15 loops=1)
2   -> Nested loop inner join (cost=359.20 rows=15) (actual time
   =5.603..5.822 rows=15 loops=1)
3     -> Index scan on J using idx_job_maxsalary (reverse) (cost=1.70
   rows=15) (actual time=0.087..0.177 rows=19 loops=1)
4     -> Filter: <in_optimizer>((C.Industry,J.MaxSalary),(C.Industry,J.
   MaxSalary) in (select #2)) (cost=0.25 rows=1) (actual time
   =0.297..0.297 rows=1 loops=19)
5       -> Single-row index lookup on C using PRIMARY (CompanyName=J.
   CompanyName) (cost=0.25 rows=1) (actual time=0.004..0.004
   rows=1 loops=19)
6       -> Select #2 (subquery in condition; run only once)
7         -> Materialize with deduplication (actual time
   =5.477..5.477 rows=101 loops=1)
8           -> Aggregate using temporary table (actual time
   =5.402..5.402 rows=101 loops=1)
9             -> Nested loop inner join (cost=800.70 rows=1424)
   (actual time=0.041..4.038 rows=1514 loops=1)
10               -> Covering index scan on Job using
   idx_job_companyname_maxsalary (cost=302.30
   rows=1424) (actual time=0.032..0.495 rows
   =1514 loops=1)
11               -> Single-row index lookup on Company using
   PRIMARY (CompanyName=Job.CompanyName) (cost
   =0.25 rows=1) (actual time=0.002..0.002
   rows=1 loops=1514)
```

Analysis Adding the composite index further optimizes the subquery by reducing I/O operations. The covering index scan on Job using idx_job_companyname_maxsalary improves the performance of the subquery. While the cost metric did not change, the actual execution time decreased, indicating improved performance.

3.3.3 Indexing Design 3: Index on Company.Industry

Index Creation

```
1 CREATE INDEX idx_company_industry ON Company(Industry);
```

Analysis Indexing `Company.Industry` may help optimize the equality comparison in the subquery. However, given the subquery's structure, the materialization step may still be a bottleneck. The impact may not be significant for the current dataset size.

3.4 Final Index Design and Justification

Final Indexes Selected:

1. Index on `Job.MaxSalary`
2. Composite Index on `Job(CompanyName, MaxSalary)`

Justification The combination of these indexes reduced the query cost from 800.70 to 359.20 and improved the overall execution performance. The index on `Job.MaxSalary` allows efficient retrieval of the highest salaries, while the composite index optimizes the subquery by enabling covering index scans. These indexes address the primary bottlenecks in both the main query and the subquery without negative impacts.

4 Advanced Query 4: Companies with Most Job Openings

4.1 Query Overview

Objective: Select the companies with the most job openings.

Listing 5: Advanced Query 4

```
1 SELECT C.CompanyName, COUNT(J.JobId) AS TotalJobs
2 FROM Company C
3 JOIN Job J ON C.CompanyName = J.CompanyName
4 GROUP BY C.CompanyName
5 ORDER BY TotalJobs DESC
6 LIMIT 15;
```

4.2 Baseline Performance Without Additional Indexes

4.2.1 EXPLAIN ANALYZE Output Before Indexing

```
1 -> Limit: 15 row(s) (actual time=6.929..6.931 rows=15 loops=1)
2   -> Sort: TotalJobs DESC, limit input to 15 row(s) per chunk (actual
3     time=6.928..6.929 rows=15 loops=1)
4     -> Table scan on <temporary> (actual time=6.666..6.807 rows=1044
5       loops=1)
6       -> Aggregate using temporary table (actual time=6.662..6.662
7         rows=1044 loops=1)
7         -> Nested loop inner join (cost=800.70 rows=1424) (actual
            time=0.124..4.977 rows=1514 loops=1)
            -> Covering index scan on J using
                idx_job_companyname_maxsalary (cost=302.30 rows
                    =1424) (actual time=0.089..0.684 rows=1514 loops=1)
            -> Single-row covering index lookup on C using PRIMARY
                (CompanyName=J.CompanyName) (cost=0.25 rows=1) (
                    actual time=0.003..0.003 rows=1 loops=1514)
```

4.2.2 Analysis

- **Total Cost Before Indexing:** 800.70
- **Observation:**
 - The query performs a covering index scan on Job using an existing index.
 - The join between Company and Job uses the primary key and foreign key on CompanyName.
- **Bottleneck Identified:** The main cost comes from the nested loop join over all Job records and the grouping operation.

4.3 Indexing Strategies

4.3.1 Indexing Design 1: Composite Index on Job(CompanyName, JobId)

Index Creation

```
1 CREATE INDEX idx_job_companyname_jobid ON Job(CompanyName, JobId);
```

EXPLAIN ANALYZE Output After Adding Index

```
1 -> Limit: 15 row(s) (actual time=5.811..5.813 rows=15 loops=1)
2   -> Sort: TotalJobs DESC, limit input to 15 row(s) per chunk (actual
3     time=5.810..5.811 rows=15 loops=1)
4     -> Table scan on <temporary> (actual time=5.525..5.666 rows=1044
5       loops=1)
6       -> Aggregate using temporary table (actual time=5.521..5.521
7         rows=1044 loops=1)
          -> Nested loop inner join (cost=800.70 rows=1424) (actual
            time=0.107..3.995 rows=1514 loops=1)
              -> Covering index scan on J using
                idx_job_companyname_jobid (cost=302.30 rows=1424) (
                  actual time=0.073..0.506 rows=1514 loops=1)
              -> Single-row covering index lookup on C using PRIMARY
                (CompanyName=J.CompanyName) (cost=0.25 rows=1) (
                  actual time=0.002..0.002 rows=1 loops=1514)
```

Analysis The composite index allows for a covering index scan on Job, reducing disk I/O and improving cache utilization. The actual execution time decreased from 6.929 seconds to 5.813 seconds. Although the cost metric remains the same, the performance gain is evident from the reduced execution time.

4.3.2 Indexing Design 2: Index on Job(CompanyName)

Index Creation

```
1 CREATE INDEX idx_job_companyname ON Job(CompanyName);
```

EXPLAIN ANALYZE Output After Adding Index No significant change in cost from baseline performance observed; the total cost remains at 800.70.

Analysis The query execution with the index on Job(CompanyName) shows a total cost of 800.70, which remains unchanged from the baseline performance prior to indexing. The execution time for retrieving the top 15 companies with the most job openings is 9.605 seconds. While the index facilitates faster lookups during the join between the ‘Company’ and ‘Job’ tables, the primary bottlenecks persist in the aggregate operation and the temporary table scan, which processed 1044 rows and took 9.344 to 9.483 seconds. The nested loop inner join continues to efficiently utilize the covering index scan on Job, with minimal impact on the overall performance. Consequently, despite the improved efficiency of individual lookups, the unchanged cost structure indicates that further indexing strategies may be necessary to enhance performance, particularly regarding the aggregation and temporary table management.

Index Removal

```
1 DROP INDEX idx_job_companyname ON Job;
```

4.3.3 Indexing Design 3: Index on Job(JobRole)**Index Creation**

```
1 CREATE INDEX idx_job_jobrole ON Job(JobRole);
```

EXPLAIN ANALYZE Output After Adding Index No significant change in cost from baseline performance observed; the total cost remains at 800.70.

Analysis The execution time for the query using the index on `Job(JobRole)` shows an actual time of 5.162 seconds, which indicates a slight improvement in performance compared to previous executions. However, the cost metrics remain unchanged at 800.70, suggesting that the addition of this index did not impact the overall cost of the join and aggregation operations. The query still utilizes the existing covering index on `Job(idx_job_companyname_jobid)` effectively, as

Index Removal

```
1 DROP INDEX idx_job_jobrole ON Job;
```

4.4 Final Index Design and Justification

Final Index Selected: Composite Index on `Job(CompanyName, JobId)`

Justification The composite index improves query performance by enabling index-only scans, reducing I/O operations, and optimizing the join and aggregation. There are no adverse effects on other queries or write operations, making it a suitable indexing strategy for this query.

- Advanced Query 1 and 2: Full-text indexing on Job.JobRole significantly reduced query costs by effectively handling text searches.
- Advanced Query 3: Indexes on Job.MaxSalary and Job(CompanyName, MaxSalary) addressed sorting and joining bottlenecks.
- Advanced Query 4: A composite index on Job(CompanyName, JobId) improved performance through efficient joins and reduced I/O.